

Logging

Logging

- Process of recording an events that happened in application in a specific point in time
- Gives us understanding of what is software behavior
- Often has low priority in development
- Logging != Monitoring
- Process of repeatedly collecting data metrics of application in a specific point in time

Source

- Application
 - exceptions, info
- Server
 - web servers access logs, database long query logs
- System
 - boot log, kernel logs, cron job logs

Destination

- Files
 - often a default, can become big in size, problem to scale
- API
 - easy to configure, external dependency, expensive
- Database
 - not common, slow to write in db, increase db size
- Stdout
 - responsibility of the environment, docker as an example

Format

- Plain-text
 - free-form text, most common, formatted
- Structured
 - more advanced, JSON, easy to search
- Binary
 - systemd, mysql binlogs, not human readable

Decentralized vs Centralized

- Decentralized
 - each server has logs, hard to access when scaled
- Centralized
 - logs in one place, easy to access

Cloud-based vs Self-hosted

- Cloud-based
 - centralized server on a cloud as SaaS, easy to setup and maintain
 - loggly, papertrail, splunk
- Self-hosted
 - build to fit our needs, managed by us
 - harder to setup, lower monthly costs, easier to scale

Log levels

- Info
- Warn
- Error
- Fatal
- Debug
- Trace
- 1-10

Log rotation

- Important when logging into files
- Process of compressing and archiving the log files
- Rotate on file size or time interval
- Archived files are deleted after a period of time
- Without it, big files can crash the server

Log alerts

- Sending alerts based on log entries

Components

- Log collectors
 - collect logs from different sources, low affect on server performance
 - scribe, flume, logstash, fluend, graylog2
- Storage / Search engine
 - store logs and perform very fast searches on large data
 - elasticsearch, solr
- Visualization dashboards
 - view, search and visualize logs
 - kibana, graylog, grafana

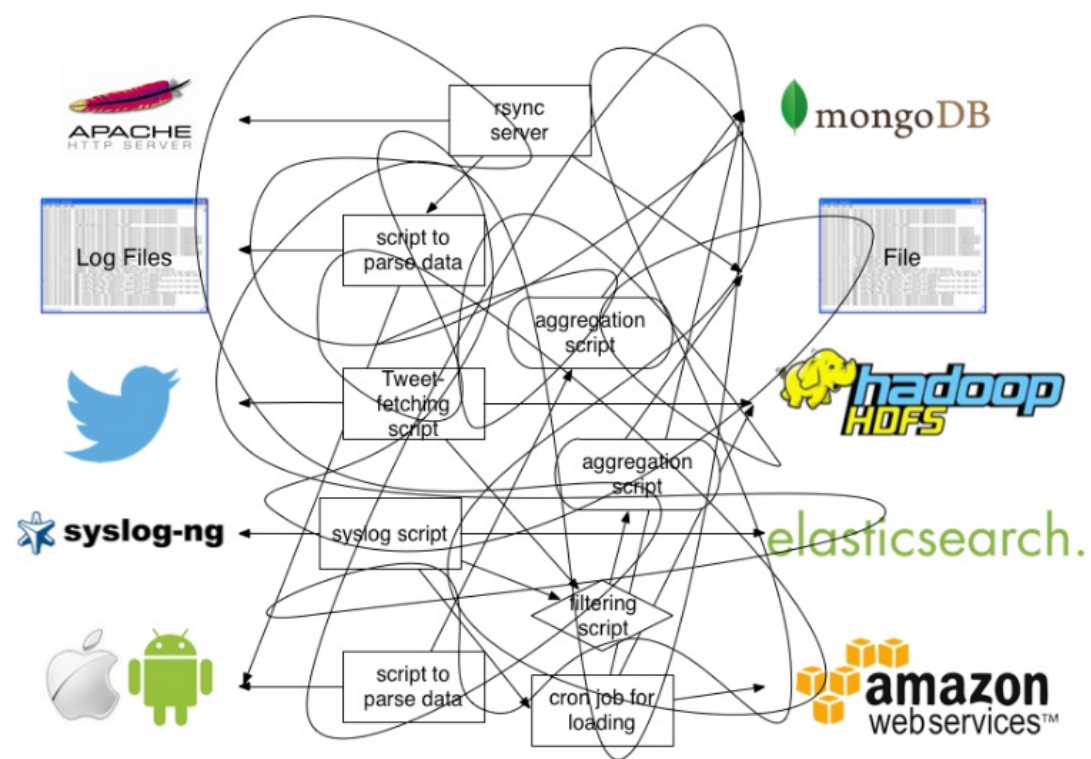
Stack

- Graylog
 - covers all components, can use other log collectors
- ELK
 - elasticsearch, logstash, kibana
 - battle-tested, standard in centralized logging
- EFK
 - elasticsearch, fluentd, kibana

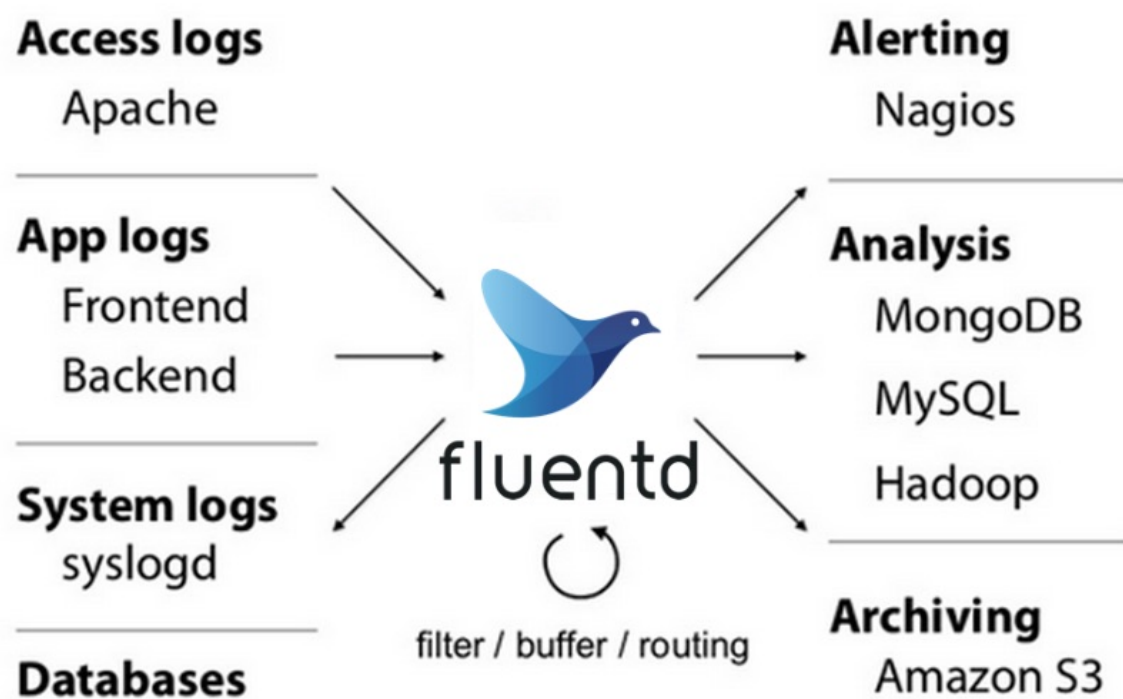
Logstash vs Fluentd

- Logstash
 - used to consume more memory
 - event routing based on algorithmic statements
 - centralized plugin repository
 - recommended for traditional systems with VMs
- Fluentd
 - event routing based on tags
 - decentralized plugin repository
 - built by cloud native computing foundation

Fluend



Before Fluentd



After Fluentd

Fluend configuration

- @include

@include /path/to/config.conf

- Source

```
<source>  
  @type http  
  port 9880  
</source>
```

- Match

```
<match myapp.tag>  
  @type file  
  path /var/log/fluent/myapp.tag.log  
</match>
```

Fluend configuration

- Filter

```
<filter myapp.access>  
  @type record_transformer  
  <record>  
    host_param "#{Socket.gethostname}"  
  </record>  
</filter>
```

- System

```
<system>  
  log_level error  
</system>
```


Fluend event

- Apache access log record

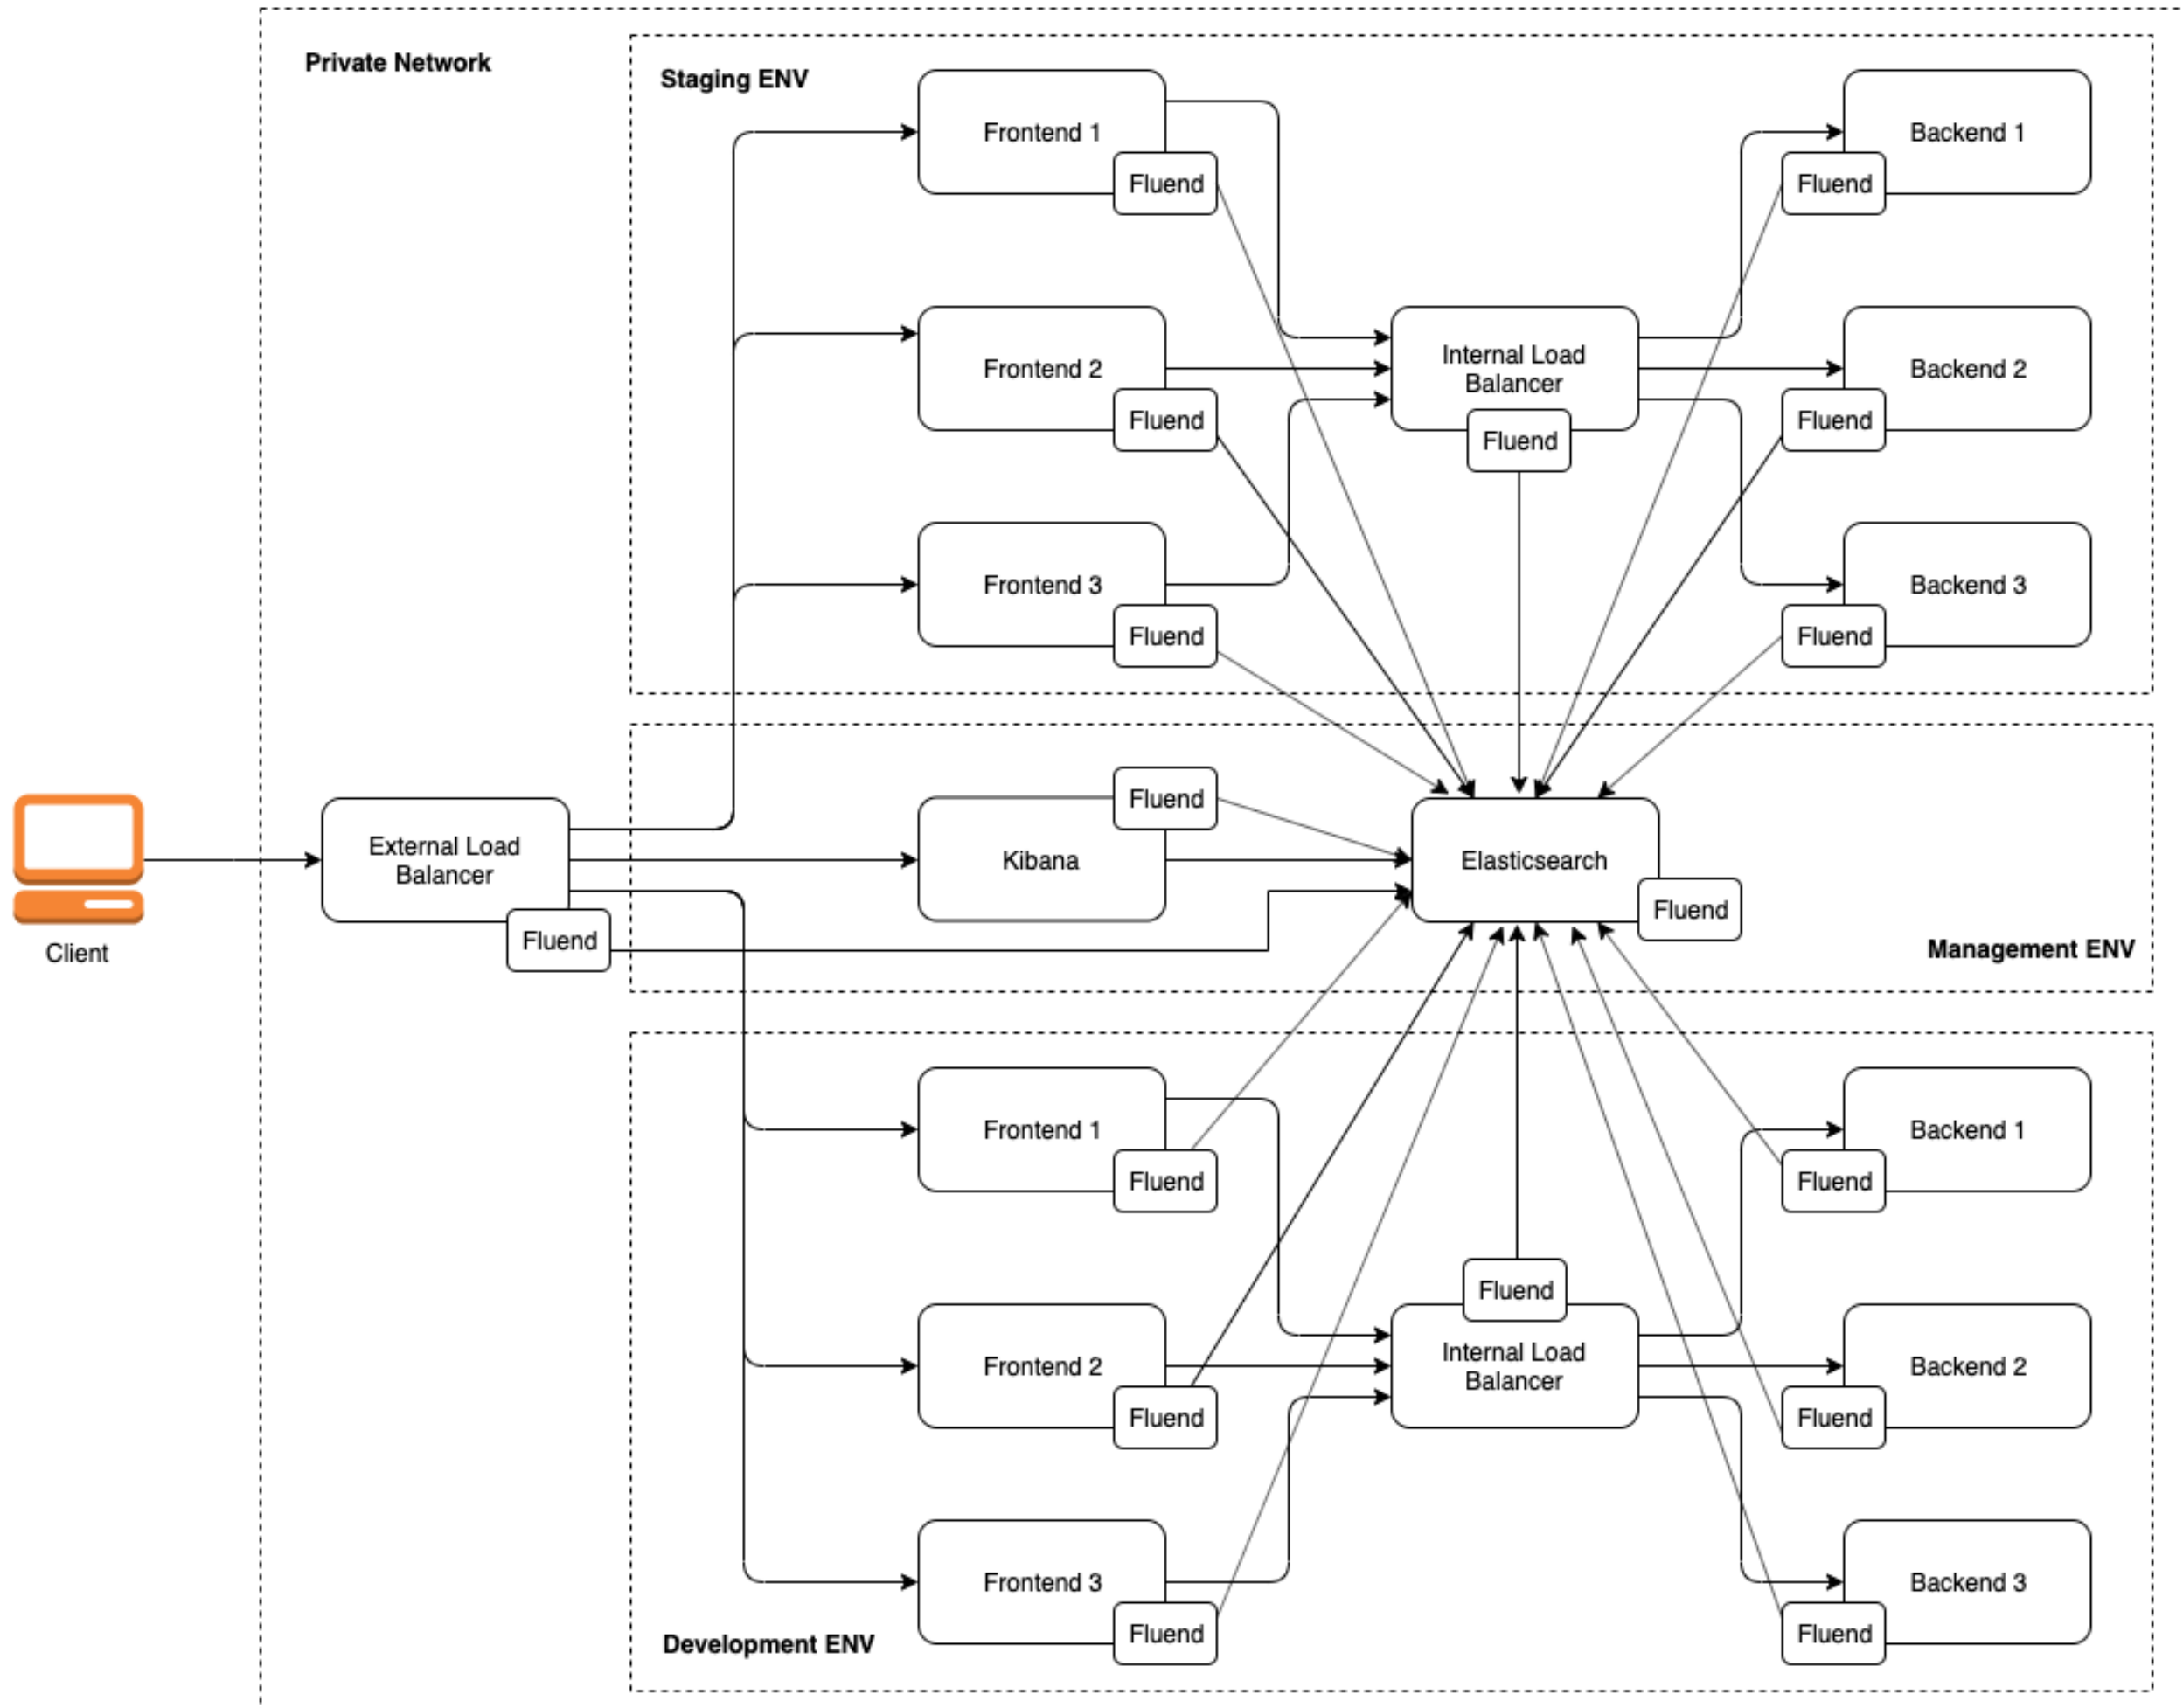
192.168.0.1 - - [28/Feb/2013:12:00:00 +0900] "GET / HTTP/1.1" 200 777

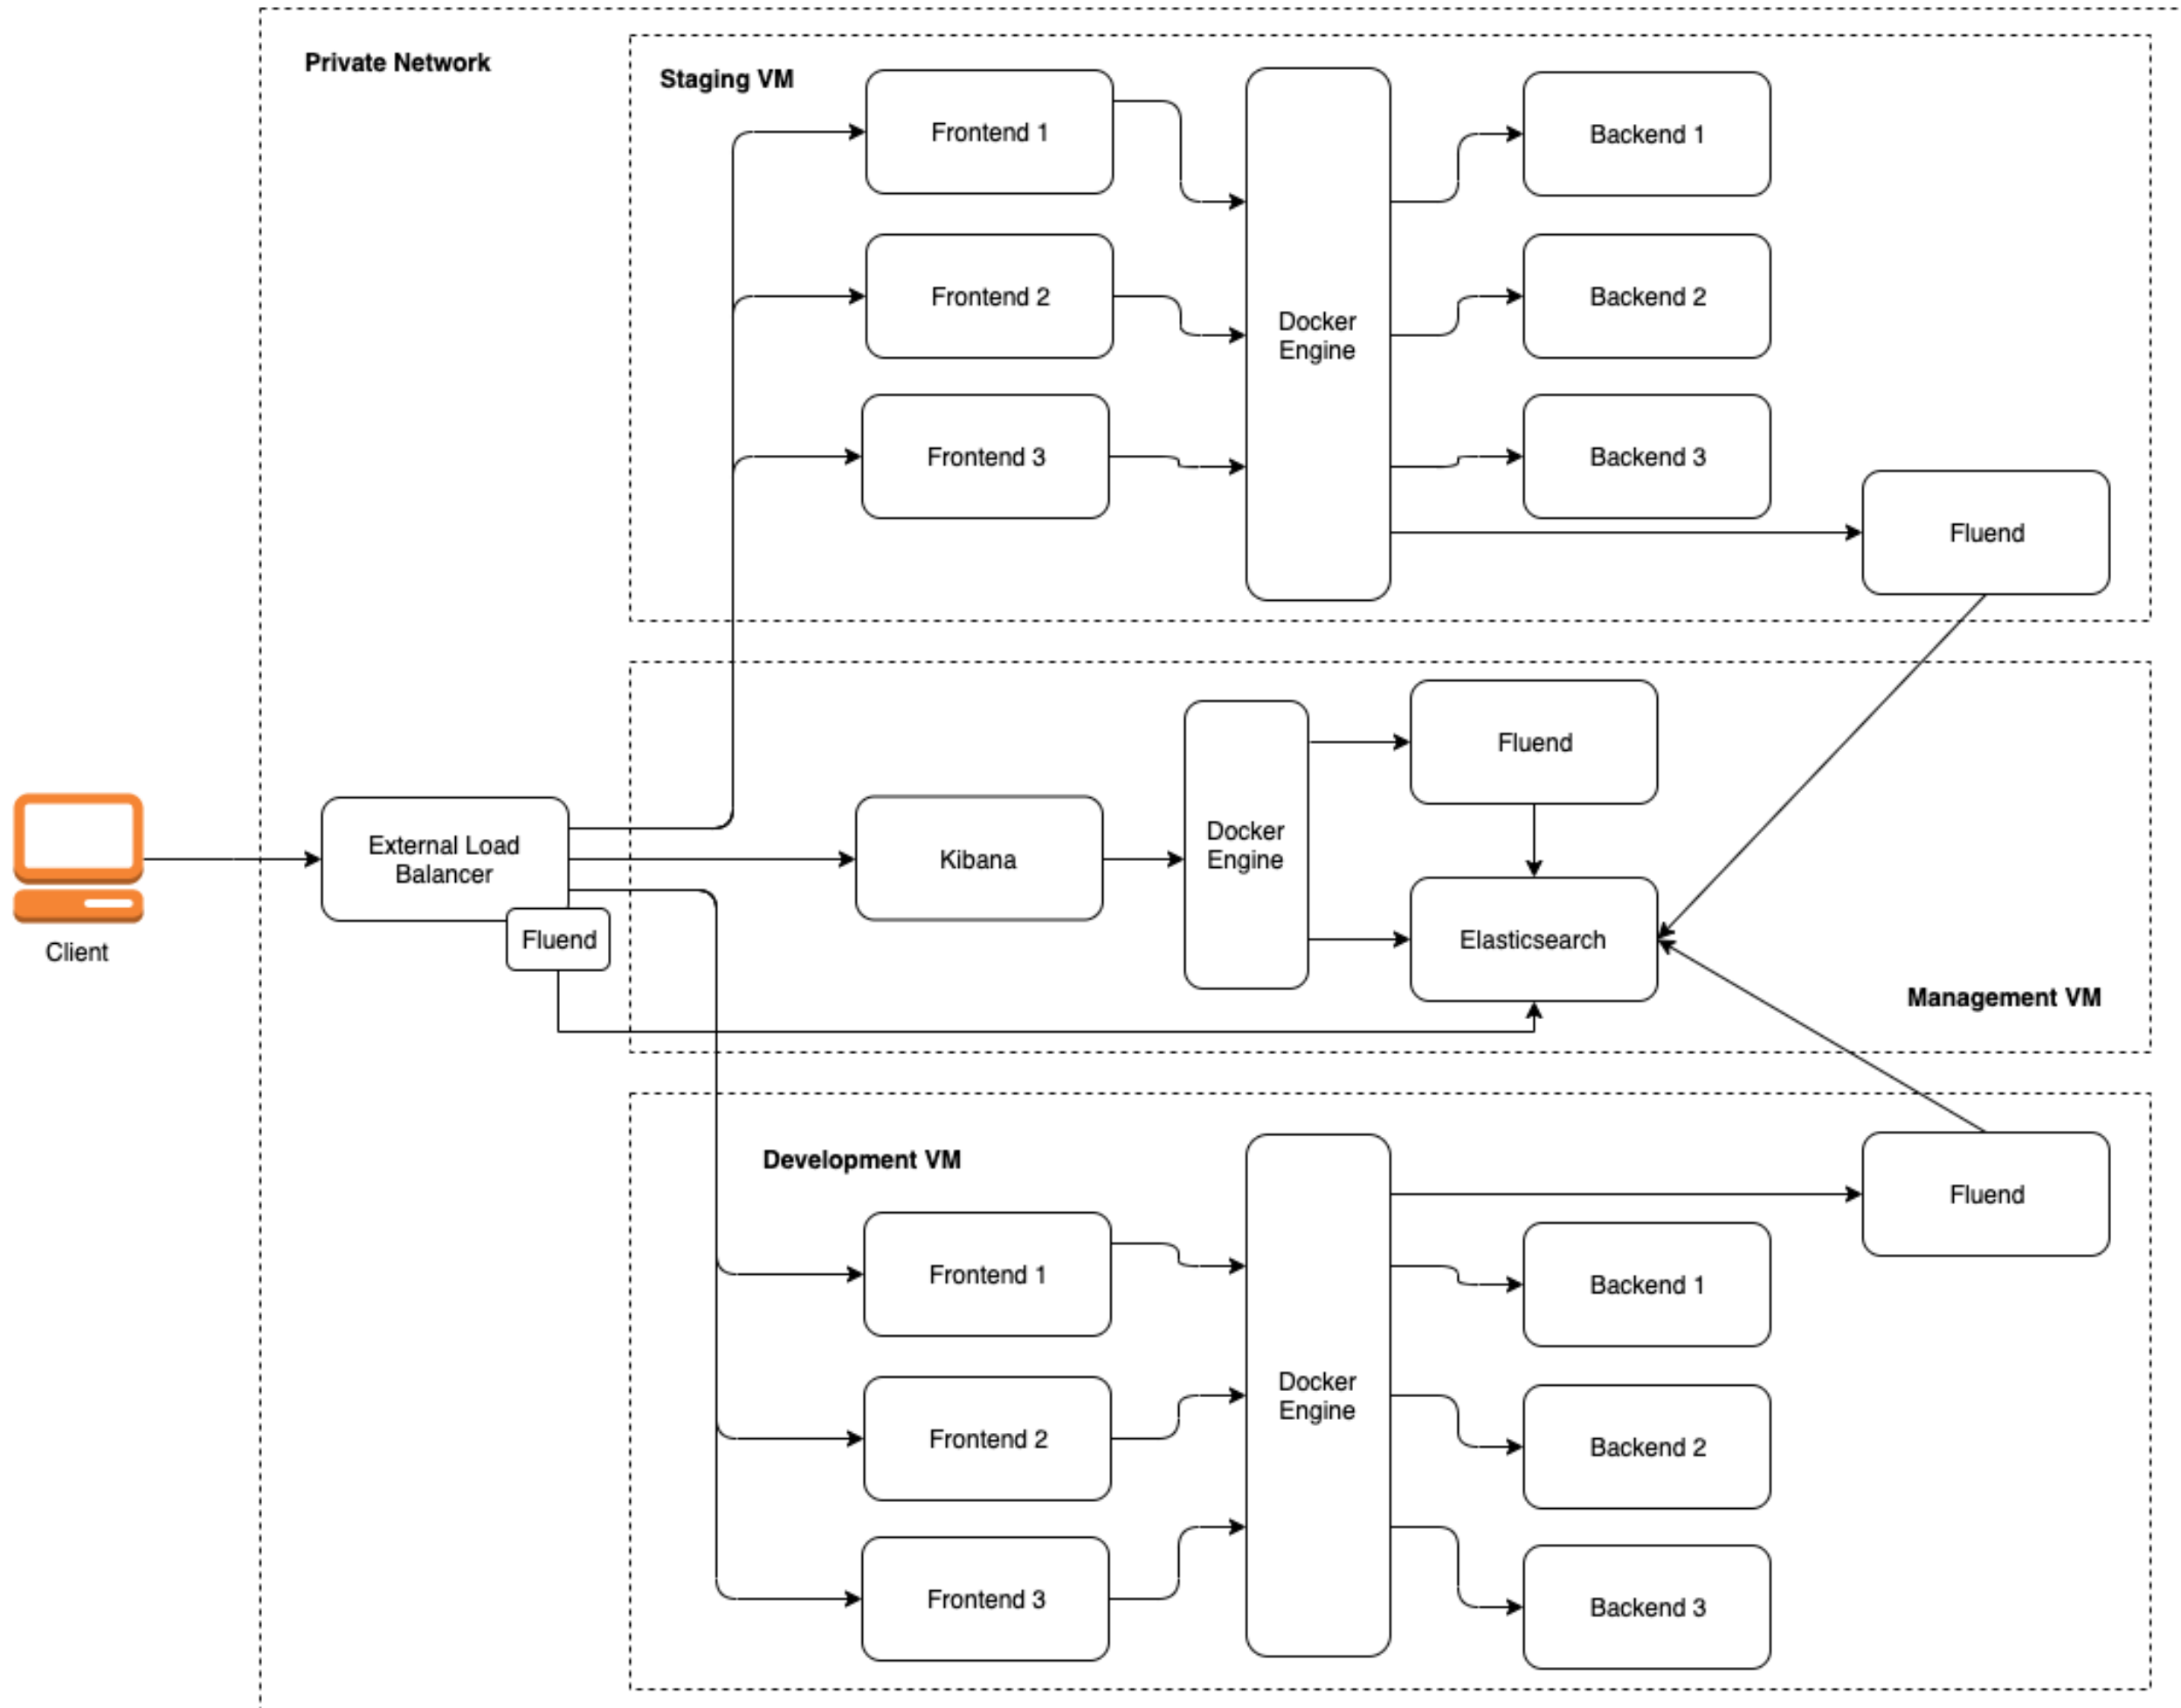
- Fluend event

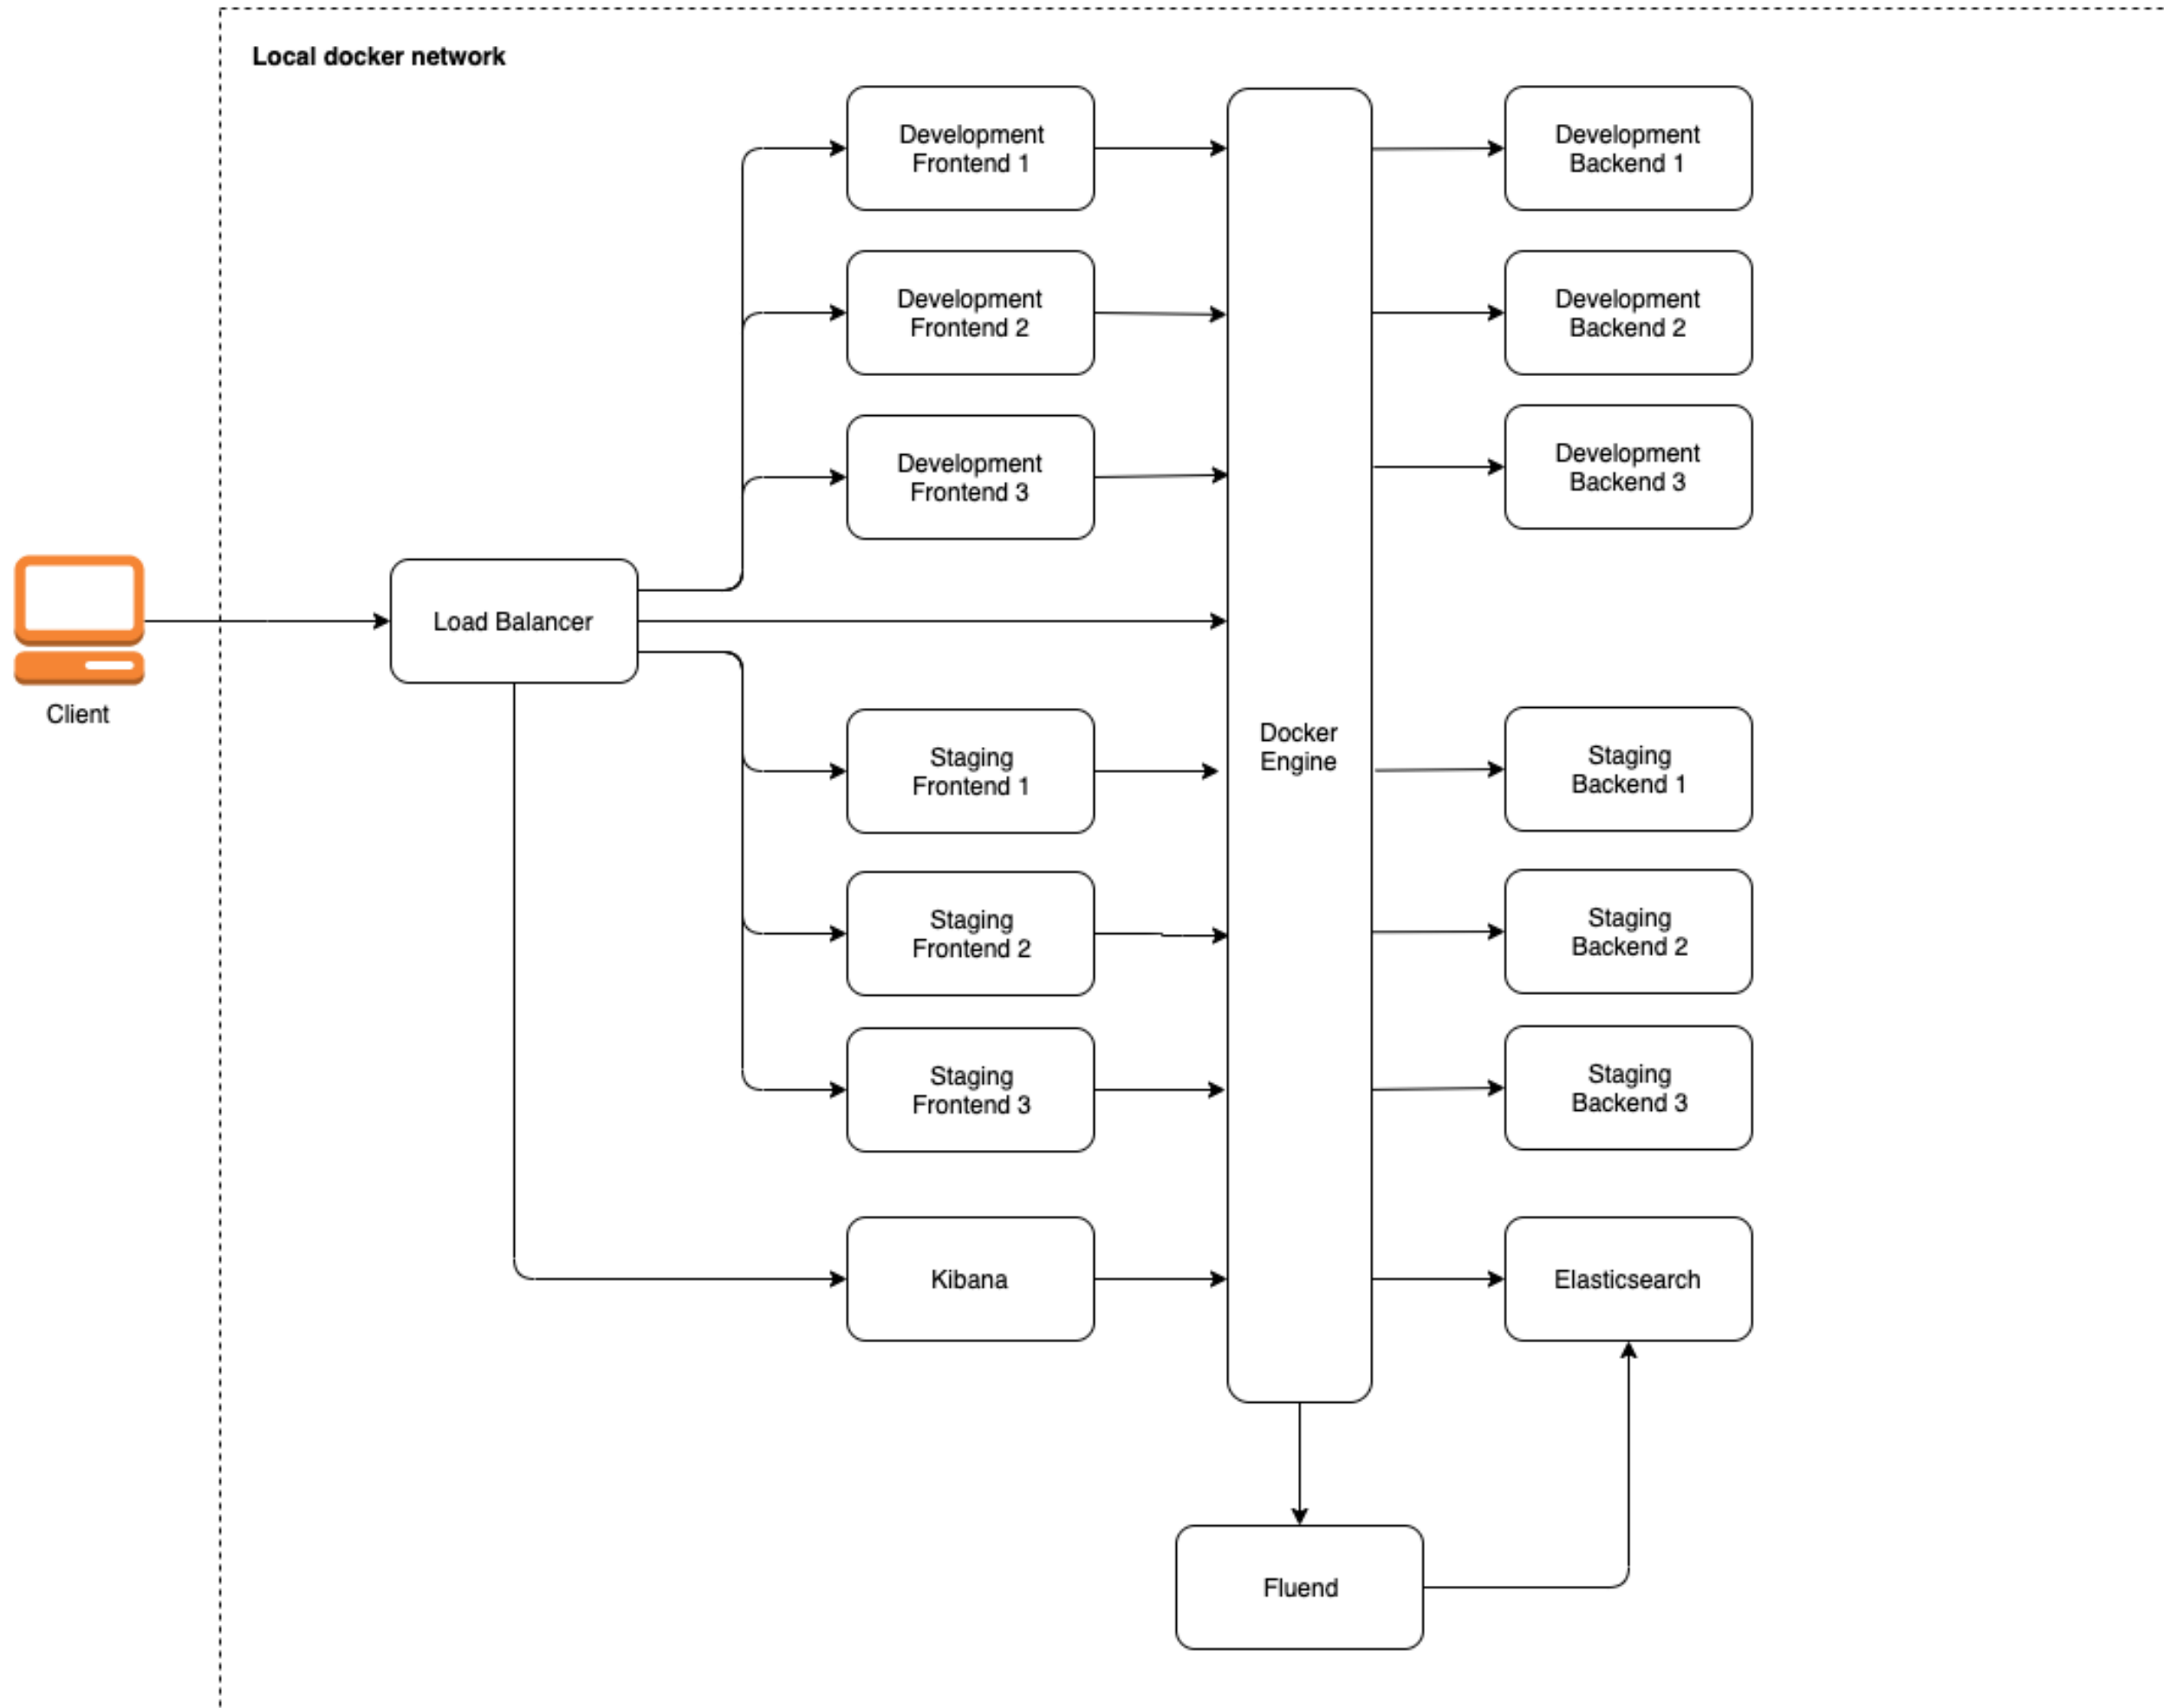
tag: apache.access # set by configuration

time: 1362020400 # 28/Feb/2013:12:00:00 +0900

record: {"user": "-", "method": "GET", "code": 200, "size": 777, "host": "192.168.0.1", "path": "/"}







Demo

<https://github.com/aleksmark/centralized-demo>

Resources

- <https://hackernoon.com/part-1-building-a-centralized-logging-application-5a537033da0a>
- <https://www.loomsystems.com/blog/single-post/2017/01/30/a-comparison-of-fluentd-vs-logstash-log-collector>
- <https://codefarm.me/2018/06/29/elasticsearch-fluentd-kibana-docker-compose/>
- <https://docs.fluentd.org/quickstart/life-of-a-fluentd-event>
- <https://docs.docker.com>
- <https://docs.fluentd.org>